

# Variation on a Scheme: multiple implicit first-class continuations and a general attempt at understanding continuations

Jonathan Lam

August 23, 2021

## 1 Introduction

This is for the completion of ECE491: Structure and Implementation of Computer Programs, based on the text of the same name by Abelson and Sussman. The assignment was to implement an extension of the interpreter presented in sections 4.1-4.4.

This work extends the implementation presented in section 4.3: “Variations on a Scheme – Nondeterministic Computing.” The section uses implicit continuations to implement the `amb` keyword. This work extends that in two ways by:

- providing an arbitrary number of continuations, rather than only two (success and fail); and
- exposing continuations as first-class objects using the `call/cc` interface

Additionally, a significant portion of this project was dedicated to exploring the theory and use cases of continuations, so there will be some backgrounds and examples of common use cases.

## 2 Background

Simply due to curiosity, this project became largely exploratory in nature. This project ultimately aims to answer the question: “What are continuations and when should we use them?” Thus this background section forms the bulk of the report and the research effort, despite the original goal to implement a new feature.

From Matt Might’s blog post *Continuations by example: Exceptions, time-traveling search, generators, threads, and coroutines*:

Continuations are the least understood of all control-flow constructs. This lack of understanding (or awareness) is unfortunate, given that

continuations permit the programmer to implement powerful language features and algorithms, including exceptions, backtracking search, threads, generators and coroutines.

I think part of the problem with continuations is that they're always explained with quasi-metaphysical phrases: "time travel," "parallel universes," "the future of the computation." I wrote this article so that my advanced compilers students could piece together how continuations worked by example.

I wholly agree with this sentiment. There are a dozen different ways to interpret continuations, and each interpretation lends its own insights. In this section I will attempt to give a better idea of what continuations are, what they are capable of, how they are used, and even how they may be implemented – all this just to get an inkling of what this mysterious control structure is.

## 2.1 What are continuations?

As Might points out, a common way to describe continuations are that they are the "future" of an expression. Luckily, Scheme's syntax is able to provide an objective way to view the future of an expression: the future of an expression is the containing expression. This is due to the fact that parameters are fully evaluated before the function is applied. Consider Figure 1. The result of the subexpression `(* a b)` is the addition by `c`, so we can express its continuation as a procedure that performs this action: `(lambda (val) (+ val c))`. The futures (continuations) for each subexpression are also shown in the figure.

Another way to think about it is that the continuation for any expression is the (implicit) callback for a given computation. After `(* a b)` is fully computed, then it should send its result to the lambda function representing its continuation. This should not feel too strange to those familiar with asynchronous programming: often the result of an asynchronous computation is passed to a specified callback or handler function.

Continuations become very useful if the programmer has access to them. For example, since the `(* a b)` is an ordinary procedure, we can save this procedure to a variable and call it like an ordinary function. If the programmer saves the continuation of that subexpression to a variable `cont`, then it can be invoked. `(cont 42)` with 5 assigned to `c` will produce 47.

Of course, this is not a complete characterization, and the rest of the background section will be used to foster a more complete understanding of continuations. At many times, continuations will feel like an analogous construct in another language; it may seem like a "first-class return" or a "snapshot of the program state"; no single characterization is universally better than another.

Note that each expression or statement (in any programming language) has an implicit continuation, and continuations are an incredibly general concept. For example, the continuation of a statement in a sequence of statements is the next statement. However, we are only concerned with *reified* continuations in

```

;;; consider the following expression: a*b+c
(+ (* a b) c)

;;; future of a:      (lambda (val) (* val b))
;;; future of b:      (lambda (val) (* a val))
;;; future of (* a b): (lambda (val) (+ val c))
;;; future of c:      (lambda (val) (+ (* a b) val))

```

Figure 1: Continuation as the future of an expression

this report (i.e., continuations with a concrete implementation in the programming language), and often *first-class* reified continuations (continuations that appear as first-class objects to the programmer).

### 2.1.1 The `call/cc` API

In the previous section, continuations are characterized as procedures; but how do we get these procedures? In Scheme, implicit continuations are exposed to the programmer using the `call-with-current-continuation` procedure, which is often aliased to `call/cc`. `call/cc` takes as argument a procedure that takes the current continuation as a parameter. The continuation is a first-class object – it can be stored in variables, passed as a parameter to functions, returned from functions, and be a member of complex data structures.

`call/cc` will then execute the procedure. If the continuation is not invoked, then the return value of the procedure is passed to the continuation of the `call/cc` expression. If and when the continuation is invoked, then the value given during the invocation is passed to the continuation of the `call/cc` expression. Several of `call/cc` usage are shown in Figure 2.

### 2.1.2 Continuation-passing style (CPS)

Continuation-passing style is a very instructive tool for discussing continuations at a functional level.

First-class continuations can be implemented in any language with lambdas (functions with lexical closures), by performing a syntactical transformation of functions to CPS, which is shown in Figure 3. Looking at the CPS form, we can characterize a continuation in the following ways:

- Each function takes its continuation as an extra parameter.
- A continuation is only invoked as a tail-call (if a continuation is called, its “then-current continuation will be abandoned”, so a non-tail call would be wasteful).
- A continuation takes the result of the previous instruction as its singular parameter.

```

;; the original expression
(+ 1 (+ 2 (+ 3)))
;; => 6

;; continuations are optionally-resumable:
;; i.e., may not be invoked at all
(+ 1 (call/cc (lambda (cc) (+ 2 (+ 3)))))
;; => 6

;; invoking a continuation causes the current
;; continuation to be discarded; the continuation
;; of `(cc 3)` is `(lambda (val) (+ 2 val))`, and
;; this is never invoked
(+ 1 (call/cc (lambda (cc) (+ 2 (cc 3)))))
;; => 4

;; continuations are multiply-resumable:
;; i.e., may be invoked multiple times
(define cont '())
(+ 1 (call/cc (lambda (cc) (set! cont cc) (+ 2 (+ 3)))))
(cont 4) ;; => 5
(cont 5) ;; => 6

```

Figure 2: Sample usage of call/cc

```

;;; without continuations; uses regular return
(define (add a b)
  (+ a b))
(define (mult a b)
  (* a b))
(define (a*b+c a b c)
  (add (mult a b) c))

(display (a*b+c 1 2 3)) ;; => 5

;;; continuations using CPS
(define (add a b cont)
  (cont (+ a b)))
(define (mult a b cont)
  (cont (* a b)))
(define (a*b+c a b c cont)
  (mult a b (lambda (a*b)
              (add a*b c cont))))

(a*b+c 1 2 3 display) ;; => 5

```

Figure 3: Continuation passing style example

- A function written using CPS *never returns normally* – it must exit by (tail-)calling a continuation.

Note that continuations in CPS are ordinary functions (“callbacks” in event-driven coding), and thus can be used to implement continuations in a language that doesn’t support implicit continuations. CPS is used by Abelson and Sussman to implement continuations in their nondeterministic interpreter, and this is extended for this project.

Other ways to implement continuations will be discussed in §2.5.2.

## 2.2 Nondeterministic computing and the `amb` keyword

Abelson and Sussman implement the `amb` keyword in Scheme. This keyword takes a set of possible values as input, and produces a value<sup>1</sup> that satisfies all of the assertions placed on the value in the future. For example, the example shown in Figure 4 may assign to `a` either 3 or 5.

This can be implemented using a search over all the possible solutions. What is really impressive is that we are able to state SAT problem in a declarative

<sup>1</sup>It is “nondeterministic” because the keyword only has to return *a* value that satisfies the assertions (conditions). The exact choice of value is not important. This is a SAT-solver, as Matt Might implements using CPS in his blog post on continuations.

```
(define a (amb '(1 2 3 4 5)))  
(require a odd?)  
(require a prime?)
```

Figure 4: Sample usage of `amb`

manner, rather than embedding it in an algorithm<sup>2</sup>. Even more impressively, we are able to state the constraints on an ambiguous value in its *future*, and possibly after the value has already been used.

The key part to this implementation is the implicit<sup>3</sup> use of continuations. In this case, an additional error continuation is provided to handle the case of a condition failing – this causes the next value to be tried (if any). This also necessarily undoes any side effects (such as variable mutation) to effective “time travel” to the point in the program where the ambiguous value is declared.

Functionally, this can be thought of as a try/catch statement, without the user having to explicitly code the control flow (and simply because there is not a builtin try/catch statement in Scheme).

The interpreter from section 4.3 is based off of the version in 4.1, which uses a procedure-based intermediate representation (IR) for all expressions. The expression is first parsed using the Scheme `read` procedure, and then semi-compiled into this IR to avoid parsing every time the expression is encountered (if the expression is invoked multiple times).

The form of the procedure representing the compiled expression (the expression IR) is shown in Figure 5. Whenever the expression is invoked, this lambda is called with the current runtime environment.

The analogous expression IR for the nondeterministic backtracking interpreter from section 4.3 is shown in Figure 6. At runtime, each expression is also given two continuations as parameters – this is known as continuation-passing style (CPS). These continuations are ordinary procedures that should be used to pass around the results of computations. The success continuation takes the value from the previous computation and the failure continuation. The failure computation is special, its only purpose being to discard the value and reverse the stack and side effects so that the next value can be tried – it does not need the value from the previous computation nor the success continuation.

<sup>2</sup>The same is true for generators, and is what makes them so powerful.

<sup>3</sup>This is my own terminology. I use “implicit” to refer to (reified) continuations that are baked into the language and can be used by the interpreter to perform control flow. Optionally, the interpreter may choose to expose these continuations via interfaces such as `call/cc`. I use “explicit” to refer to CPS, in which reified continuations are implemented by the user as explicit procedure calls.

```
(lambda (env)
  ;; env stores the symbol table for the current scope
  ...)
```

Figure 5: Expression IR from section 4.1

```
(lambda (env succeed-cont fail-cont)
  ;; succeed-cont is of the form (lambda (value fail-cont) ...)
  ;; fail-cont is of the form (lambda () ...)
  ...)
```

Figure 6: Expression IR from section 4.3

## 2.3 Continuations versus other control-flow constructs

This section is structured similarly to the Wikipedia page on coroutines, which is an excellent resource on coroutines.

### 2.3.1 Continuations vs. `gotos`

Gotos are the software equivalent to a jump instruction. They are very simple and limited. For example, they can only jump to other locations in the current function; jumping to another function without modifying the stack would cause relative-addressing (local variables) to break. Also unlike `gotos`, continuations usually return a value.

Continuations are more similar to C's `setjmp` and `longjmp` functions, which essentially save and restore a history of the stack. Not only does this allow you to jump to different functions, it also allows you to jump back multiple levels of the stack. However, reified continuations are more powerful (and expensive) in that they are both optionally-resumable and multiply-resumable<sup>4</sup>.

In his blog post about continuations, Matt Might suggests an idiom that behaves very similarly to `setjmp/longjmp`. For those who are unfamiliar with these functions, `setjmp` is similar to `fork` in that it returns different values depending on the context. `setjmp` is encountered either it is the next statement to execute, or when `longjmp` jumps to it. In the first case, it will return a falsy value, and in the latter it will return the value specified in the `longjmp` call.

Might does not explicitly mention `setjmp/longjmp` in his blog post, but the interface is remarkably similar. See a simple error-handling program with a nonlocal jump in C/C++ is shown in Figure 7. The equivalent program in Scheme, using the `setjmp` idiom (Might calls this `current-continuation` in his examples), is shown in Figure 8.

<sup>4</sup>In order to be multiply-resumable, a continuation essentially needs to create a copy of the stack whenever it is executed. See the section on implementation

```

main() {
    jmp_buf env;
    int val;

    val = setjmp (env);
    if (val) {
        fprintf (stderr, "Error %d happened", val);
        exit (val);
    }

    /* code here */

    longjmp (env, 101); /* signaling an error */
}

```

Figure 7: Sample usage of `setjmp/longjmp`. Source: [cplusplus.com](http://cplusplus.com): `setjmp`

Note that both `goto` and `setjmp` may be dangerous if there are side effects in the jumped code. `goto` is very dangerous since it can actually branch forward in a function, skipping variable initializations completely. `setjmp` is less dangerous in that it can only go back to a previous location on the stack (which should be safe), but may cause issues with dangling or incorrect pointers if cleanup code is skipped. Continuations (in Scheme) are safer because they cannot branch forward in time, and because garbage collection would manage dangling references.

### 2.3.2 Continuations vs. return

A continuation is like a return statement in that it usually passes a value and it (conceptually) unwinds the stack to a given location. However, continuations are optionally-resumable, multiply-resumable, and may unwind the stack more than one stack frame<sup>5</sup>.

### 2.3.3 Continuations vs. callbacks

In CPS, continuations *are* simply callbacks that follow a specific form. In general, first-class continuations feel very much like ordinary callbacks, in that they are invoked with the result of a computation. For example, in Javascript where event-driven coding is common, callbacks to asynchronous functions and the `Promise` API look very similar to CPS.

<sup>5</sup>The latter is useful in exception handling, or when needing to jump back many stack frames without the overhead of returning from each stack frame, such as in the backtracking nondeterministic interpreter.



```

(define (setjmp)
  (call/cc (lambda (cc) (cc cc))))

(let ([longjmp (setjmp)])
  (if [procedure? longjmp]
      ;; val == false
      (begin
        ;; code here
        (longjmp 101))
      ;; val = longjmp; val != false
      (error 'setjmp-example
              "Received error signal from longjmp"
              longjmp)))

```

Figure 8: Useful Scheme idiom using continuation that behaves like `setjmp`. Source: Continuations by example: Exceptions, time-traveling search, generators, threads, and coroutines

Of course, in the case of implicit continuations, there may be special machinery that manages program execution state that cannot be performed with a regular procedure call, but the API is the same as a procedure.

Note that callbacks are also multiply-resumable, optionally-resumable, and capture program state (capture its lexical environment). In general, we can think of continuations as a specific type of callback.

### 2.3.4 Continuations vs. monads

Haskell users may know that continuations exist as a monad (`Control.Monad.Cont`). Even for those who don't (but are familiar with monads), the continuation's structure should feel monadic:

- Continuations act like a wrapper around a computation.
- Chaining computations looks very similar to the `bind` operation.
- Continuations are a control-flow structure.

Consider Figure 9, the Haskell analog of the CPS example shown in Figure 3. Note that Haskell also has the `callCC` interface, similar to the `call/cc` procedure. Figure 10 shows a sample monad instance declaration for the continuation type. Unsurprisingly, the syntax between Scheme and Haskell turns out to be very similar, even when Scheme continuations feel like procedures and Haskell continuations feel like monads.

```

import           Control.Monad.Cont

mult             :: Int -> Int -> Cont r Int
mult a b        = return (a * b)

add             :: Int -> Int -> Cont r Int
add a b         = return (a + b)

-- a * b + c
atbpc           :: Int -> Int -> Int -> Cont r Int
atbpc a b c     = mult a b >=> add c

main = do runCont (atbpc 1 2 3) print

```

Figure 9: CPS in Haskell using `Control.Monad.Cont`

```

newtype Cont r t = Cont ((t -> r) -> r)
  -- a value of type Cont r t is of the form Cont f,
  -- where f is a function of type (t -> r) -> r

runCont (Cont f) q = f q

instance Monad Cont where
  return :: t -> Cont r t
    -- return takes a value of type t and produces a Cont
    -- computation of type t with response type r.
  return x = Cont (\q -> q x)
    -- produce a Cont computation which given a question q,
    -- applies the question to the value x.

  (>=>) :: Cont r t -> (t -> Cont r s) -> Cont r s
  x >=> f = Cont (\q -> runCont x (\v -> runCont (f v) q))

```

Figure 10: A simplified continuation monad in Haskell. Source: Cont computations as question-answering boxes

### 2.3.5 Continuations vs. threads

Continuations are like dormant, multiply-resumable threads. In particular, threads represent the current execution context of a program, and can be thought of as a snapshot of location on the stack. However, you cannot resume a thread from the same position multiple times, unlike continuations.

WikiWikiWeb’s ContinuationExplanation hints that threads can be used to implement continuations, and vice versa. It also mentions that some thread implementations, such as POSIX pthreads, are not as powerful as continuations and cannot fully implement all continuation use cases. Continuation implementation will be further discussed in §2.5.2.

Continuations and execution context (and in particular, concurrent threads of execution) is part of the recurring theme, and will appear again in the discussion of coroutines.

## 2.4 Uses for continuations

### 2.4.1 Callbacks and promises

Even in languages without implicit continuations, callbacks (that look and act a lot like continuations) are widely known and applied. Might’s blog post *By example: Continuation-passing style provides the example of a CPS-style wrapper around the legacy XMLHttpRequest Ajax API*<sup>6</sup> that may look extremely natural for Javascript users, even without knowing about continuations.

Another asynchronous programming model in Javascript is the **Promise** API, which can be thought of as a modern wrapper around callbacks with some nice features. Promises are objects that contain a computation, and allow you to specify a success and error continuation explicitly using the **Promise::then()** and **Promise::catch()** methods. Javascript promises have a number of interesting features that make them more appealing than ordinary callbacks, such as their ability to be **await**-ed to avoid “callback hell,” the ability to automatically chain continuations, and the ability to attach multiple continuations to the same operation (which would be called asynchronously due to Javascript’s event loop).

Since promises and callbacks are more or less functionally equivalent, scenarios that are expressible using callbacks are also expressible using promises. If we turn the example from Figure 3 into CPS, and then transform that to use the **Promise** API, then we would get Figure 11.

As stated earlier, continuations model only a specific subset of callbacks or promises that are useful for execution; namely a single success callback that takes a single parameter (the result of the previous operation), and (optionally) an error callback that takes a single parameter (the error message). For my project, I aim to generalize implicit continuations to the case where there are

---

<sup>6</sup>I’m not sure when this blog post was written, since the modern **fetch** API uses promises rather than callbacks.

```

add = (a, b) =>
  new Promise((succeed_cont, fail_cont) => succeed_cont(a+b));

mult = (a, b) =>
  new Promise((succeed_cont, fail_cont) => succeed_cont(a*b));

atbpc = (a, b, c) =>
  new Promise((succeed_cont, fail_cont) =>
    mult(a, b).then(apb =>
      add(apb, c).then(succeed_cont)));

atbpc(1, 2, 3).then(console.log);

```

Figure 11: CPS using Javascript **Promises**

an arbitrary number of continuations, which is similar to how a function is not limited in the number of callbacks it receives.

#### 2.4.2 Nonlocal branching

An example of nonlocal branching is the nondeterministic interpreter in *SICP* section 4.3. Whenever there is a requirement conflict, we have to unwind the stack to the location of the relevant **amb** invocation, which may be nonlocal.

A personal example I have is when I had to write an A\*-search algorithm and chose Scheme. Iterative deepening with a time limit was used; when the time limit was reached, rather than unwinding up the arbitrarily-deep search stack, a continuation served as a much simpler nonlocal goto. (Continuations were not used for the DFS.)

#### 2.4.3 Exception handling

This can be thought of as an example of nonlocal branching. To implement a try/catch block, a stack of error (failure) continuations has to be implemented alongside the ordinary (success) continuation. Matt Might shows this in his blog post about continuations.

Alternatively, rather than only passing the ordinary continuation to each function in CPS, both the success and failure continuations may be passed as parameters to every function. This may be a little bit more inefficient but has the same effect. This is how failure continuations were passed around in the nondeterministic interpreter.

#### 2.4.4 Coroutines (and generators)

As mentioned earlier, continuations are functionally very similar to (and potentially more powerful than) threads. Thus, continuations may be used to

implement cooperatively-scheduled threads<sup>7</sup>, coroutines (equivalent to cooperative threads), and generators (a subset of coroutines).

**Cooperative threading system** Continuations can be used to implement a cooperatively-scheduled threading system by maintaining a collection of continuations representing threads. Each thread may voluntarily yield to another available thread via the `yield` instruction. The collection may be a queue and the next continuation in the queue will be chosen when a thread yields; this forms a very simple round-robin scheduling algorithm. Matt Might’s blog post on continuations includes this as an example. This allows us to create concurrency (but not parallelism) in our interpreter.

Kotlin has an implementation of cooperative threading that is called coroutines (which are equivalent – see next example).

**Coroutines** A coroutine is a generalization of a subroutine (procedure) that maintains execution state between calls. It may exit either normally (by returning) or by *yielding* to another coroutine; when it is invoked again, it will resume execution from the yield point. Coroutines are functionally equivalent to cooperatively-scheduled threads.

A simple way to implement coroutines is to store a table of continuations for each function. Yielding to a coroutine would mean looking up its continuation and invoking it with the yielded value.

**Generators** Generators are a restricted type of coroutine that is very useful. A generator is a function that yields values one at a time. They are useful as iterators over arbitrary data structures. To the outsider, calling the function successively produces successive values of the collection; this works because the function saves its execution context between calls and thus doesn’t lose its position when iterating over the data structure.

Generators offer a nice solution to iteration that are roughly equivalent to streams (streams are employed in *SICP* section 4.4 as an iterator, but a generator would have worked as well). If neither streams nor generators were present, then implementing an iterator function without nonlocal state becomes difficult.

Generators are a strict subset of coroutines because the generator function can only yield to its caller, while a full coroutine can yield to any other coroutine. This simplifies the implementation and does not require a table of continuations. Rather, the two functions (generator and consumer) simply pass back and forth the new continuations and the yielded values.

Many languages implement generators due to their ubiquity, even if they do not implement coroutines.

---

<sup>7</sup>Threads are typically *preemptively scheduled*, which means that they yield control to another thread at arbitrary moments in their execution (generally to maintain performance or responsiveness). *Non-preemptive* or *cooperative* scheduling means that a thread only yields to another thread explicitly.

**Communicating sequential processes (CSP)** CSP (not to be confused with CPS) is a formal description of communication between concurrent threads of execution. Coroutines (i.e., cooperative threads) are generally the mechanism associated with CSP to allow communication between processes. A coroutine can efficiently yield its execution to a different coroutine; a normal thread complicates matters with its preemptive scheduling.

An example of this is Kotlin’s coroutines. Go’s goroutines are somewhat of a hybrid between the two: they are lightweight user-space threads that normally are preemptively scheduled, but have the ability to yield execution voluntarily. Go has a mechanism for communication between goroutines called *channels* that are intended to model CSP.

#### 2.4.5 Compiling with continuations

Continuations, and CPS in particular, have been used as a tool for expressing control flow as a compiler IR, usually alongside TCO<sup>8</sup>. This dates back to the first Scheme compiler, Rabbit, which used CPS as an IR. Since continuations replace return statements, the stack model may be replaced with a similarly efficient model involving only continuations. Tail calls can also be thought of as “gotos with parameters,” which lends itself to efficient implementation.

Two resources on this are Matt Might’s CPS Conversion and Andrew Appel’s *Compiling with Continuations*<sup>9</sup>.

### 2.5 Miscellaneous notes about continuations

#### 2.5.1 One procedure call, multiple “returns”

One property of continuations made clear by the comparison to `setjmp` with the idiom in Figure 8 is that the `call/cc` interface looks a lot like an ordinary function call to the caller, in that it will (typically<sup>10</sup>) return a value to the caller.

However, `call/cc` also creates a branch point (think `goto label`), which means that multiple values may be “returned” from that location if the continuation is invoked multiple times (recall that continuations are multiply- and optionally-resumable, just like an ordinary `goto label`).

Note that this does not mean that a function may return by calling its current continuation multiple times sequentially; the current continuation is discarded and replaced by the the invoked continuation. This is illustrated in Figure 12. The current continuation is called with value 1, and the current continuation of invoking the continuation (which is the next statement) is discarded, so control never reaches the next statement. This is equivalent to saying that invoking a continuation is only useful as a tail call, since any instructions that come after (i.e., the continuation of the invoking the continuation) are discarded.

---

<sup>8</sup>Note that CPS can be used without TCO, such as in the Chicken Scheme compiler, since it is transpiled to C (in which implementing TCO may add considerable complexity), and thus requires other optimizations to clean up the stack.

<sup>9</sup>I haven’t had the time to read these, but the latter is a famous text on the matter.

<sup>10</sup>Unless a continuation is invoked that unrolls the stack further than the calling statement.

```

(define (does-not-multiply-return)
  (call/cc (lambda (cc)
             (cc 1)      ;; expression a
             (cc 2))))   ;; expression b

;;; The continuation of expression a is expression b.
;;; Invoking the continuation `cc` discards the
;;; continuation of expression a, i.e., expression b
;;; will not be executed.

(does-not-multiply-return) ;; => 1

```

Figure 12: The current continuation gets replaced by a call to a continuation

Another way to think of this is that you cannot call `return` in C multiple times sequentially, as the second one is unreachable; however, what you can do is save the `return` statement as a first-class object and (at some later point in reachable code) invoke that statement to have the same behavior as returning from that function.

### 2.5.2 Efficient implementation

It's been hinted several times that the implementation of continuations would be similar to the implementation of threads, since both involve saving a snapshot of the execution state. Since continuations are multiply-resumable, it is more difficult to implement efficiently than a thread, since we need to maintain a pristine copy of the stack in case it is invoked more than once.

Another issue that makes it harder to optimize is that continuations have *unlimited extent* – they exist over the lifetime of the program, potentially outside of the context (closure) that they were created in. Thus it is possible to not only unroll the stack to a previous position, but to also enter a deeper position in the stack. As a result, optimizations must maintain a reference to all parts of the stack that are present at the time of the continuation being called.

*Segmented stacks* and *cactus stacks* are optimizations that allow multiply-resumable continuations unlimited extent by branching the stack and reusing memory when possible. Copy-on-write (COW) is a common memory optimization technique that also applies here to avoid copying the stack unless necessary. A more complete discussion of implementation is given in WikiWikiWeb: ContinuationImplementation.

As mentioned previously, one way to implement continuations is using CPS. CPS usually depends on tail-call optimization (TCO).

We can simplify the implementation by restricting the capability of continuations. One possible restriction is to only allow singly-resumable (but still optionally-resumable) continuations. Another restriction is to only allow up-

calls (unwinding the stack, i.e., limited extent) – these are called *escape continuations*.

### 2.5.3 Flavors of continuations

In the previous section, we have already talked about two forms of restricted continuations: single-use and escape continuations. Continuations with both of these restrictions of these are very similar to `setjmp` – this is still powerful enough for exception handling, generic nonlocal jumps. The WikiWikiWeb article on single-use continuations notes that `setjmp` can be used to implement a multi-use branch point using the very simple idiom `while(setjmp(buf));`. The Wikipedia article on continuations notes that escape continuations are a means to implement TCO.

There are also *delimited continuations*, which are a generalization of the continuations mentioned thus far. Rather than capturing the entire execution state, a part of the stack can be captured. Its API are the `shift` and `control` procedures (rather than `call/cc`).

## 2.6 Analogs in other languages

Several examples of continuations in other languages have been provided; this section serves as a summary.

In C, the flow-control capability of `goto`, `return`, and `setjmp` are all a subset of the ability of continuations. Of these, `setjmp` is the most powerful and are roughly equivalent to single-use escape continuations.

In Haskell, continuations are exposed as a monadic type. Expressing chained computations using the bind (`>>=`) operator is very natural.

Many languages have some support for generators (e.g., Python), cooperative scheduling (e.g., Kotlin’s coroutines), and CSP (e.g., Golang’s Goroutines). These may not explicitly be implemented using continuations, but the machinery is likely similar.

## 3 Multiple continuations

For this project, I decided to extend the Scheme interpreter by implementing a system of *multiple continuations* in programmer-space via a `call/cc`-like interface.

The following terminology should be established since we are working with Scheme code in the implementation of the interpreter, as well as Scheme code that may be run in the interpreter. This is non-standard – I’m not sure what the convention is for interpreter-writers.

**Programmer-space** This refers to code that the programmer may type into the interpreter.

**Interpreter-space** This refers to code in the implementation of the interpreter (including any IRs for programmer-space constructs).



### 3.1 Functional API

Functionally, multiple continuations is equivalent to CPS with an arbitrary number of continuations (callbacks), but the continuations are defined implicitly so that the programmer does not have to carry the continuations through all of the function signatures. This is a simple generalization of the single-continuation scheme and may be useful where multiple “futures” are useful, such as a default and error future in exception handling.

The implementation is based on sections 4.1 and 4.3 of *SICP*. The nondeterministic interpreter uses a CPS IR but doesn’t expose continuations to the user. It implements two continuations for every expression: a default (success) continuation, and a special-purpose backtracking (failure) continuation<sup>11</sup>. For this project, the special failure continuation is removed, and the default continuation is replaced with a list of continuations.

To the programmer, the introduction of multiple continuations change would look like the creation of a few APIs, described below, that are not present in the interpreter given in the book. The first is the familiar `call/cc` API (which does not require multiple continuations). In the case of the latter APIs, we establish the convention that the first continuation is the default continuation, the second continuation is an error continuation, and any others are user-defined.

#### 3.1.1 Exposing continuations via the `call/cc` API

This API should appear to the programmer exactly like the builtin `call/cc` implementation described in §2.1.1. In particular, it should call the provided procedure with the current (default) continuation as the sole argument. The continuation should be a first-class object, and invoking it should appear no differently than an ordinary function call. If a continuation is invoked, then the current continuation should be discarded.

This is implemented in terms of the more general case for multiple continuations; the implementation is shown in Figure 13.

#### 3.1.2 Multiple continuations via the `call/ccs` API

This API is very similar to `call/cc`, except that the argument to the function is the list of current continuations.

As an example, the code to send a value to the error continuation (to `throw` a value) is shown in Figure 13.

#### 3.1.3 User-defined continuations via the `call/new-ccs` API

This API is unlike anything in the Scheme language, so it is worth describing it here. This procedure takes two arguments: a procedure that transforms the

---

<sup>11</sup>Alternatively, the failure continuation may be implemented in an external runtime stack rather than be passed around to every function as a parameter – this may be more efficient. Exceptions may also be handled this way.

```

;;; implement `call/cc` using `call/ccs`
(define (call/cc f)
  (call/ccs (lambda (ccs)
              (f (car ccs)))))

;;; implement throw using `call/ccs`
(define (throw val)
  (call/ccs
   (lambda (ccs)
     ((cadr ccs) val))))

;;; swap the default and error continuations using `call/new-ccs`
(call/new-ccs
 ;; cc transformer function swaps first and second ccs
 (lambda (ccs)
   (cons (cadr ccs)
         (cons (car ccs)
               (caddr ccs)))))

42          ;; goes to error continuation
; (throw 42) ;; goes to default continuation
)

```

Figure 13: Intuitive of the API functions

current continuations and produces the new continuations, and an expression to invoke with the new ccs.

As an example, the code to switch the default and secondary continuations is shown in Figure 13.

## 3.2 Implementation

All of the expression syntactical analysis functions in the IR must return a CPS function, in the same manner as the nondeterministic interpreter. However, instead of receiving two continuations, it receives a list of continuations. For almost all cases, the procedure is exited by calling the first (default) continuation with the expression value.

The tricky part is that every time we invoke an IR procedure of this form, we must supply an array of continuations. If we need to modify the default continuation, then we must `cons` this to the rest of the original continuations. This makes the CPS more tedious.

To represent the continuations (ordinary procedures in interpreter-space) as procedurs in programmer-space, we must wrap them with some metadata like we do for primitive and compound procedures. In this case, we simply tag

the continuation with the symbol `'continuation`. Continuation invocation is treated the same as a primitive procedure in `execute-application`, except that the current continuations are discarded when a continuation is executed.

The implementations for `call/ccs` and `call/new-ccs` is shown in Figure 14. The implementation for `call/cc` was already shown above; it is a special case of `call/ccs`, and can be trivially written in terms of it. (Note that while the former two are *special forms*, the latter is an ordinary procedure in programmer-space.)

Note that the continuations for the user-specified continuations in `call/new-ccs` are the continuations for the `call/new-ccs` expression. This essentially automatically creates a stack of continuations. For example, throwing an exception inside a try/catch block would go to its error continuation (the catch clause). Throwing an error inside the catch clause would then go to its error continuation (which should be the error continuation of the try/catch block).

I should also add that I do not know how continuations and the `call/cc` interface are implemented in Scheme (and the implementation may differ between Schemes). This is also not intended to be a very efficient implementation, but is rather a proof-of-concept that continuations can be implemented very simply in a functional language with lambda closures, garbage collection, and TCO.

There may be a number of safety issues with `call/new-ccs`, especially with the ability to overwrite the default continuation. For this exploratory assignment, I have left it alone in all its dangerous glory, but this should probably not be productionized.

### 3.3 Examples

Implementing try/catch is very simple in programmer-space, and is shown in Figure 15 (the implementation for `throw` was already shown in Figure 13). The only caveat is that the try and catch clauses must be provided as lambdas rather than plain expressions, in order to avoid their eager evaluation. (If it were implemented as a special form, or if syntactic macros were implemented in the interpreter, plain expressions would be possible.)

Matt Might also provides an example of generators using continuations in his blog post *Continuations by example: Exceptions, time-traveling search, generators, threads, and coroutines*. While this does not use multiple continuations at all, it is a good proof-of-concept of generators, and is more evidence that my implementation of `call/cc` is correct. My implementation of the generator in programmer-space is not the most convenient form. Might writes his in terms of syntactic macros to make the syntax much more appealing. My implementation is shown in Figure 16.

### 3.4 Source code

My implementation as the files `4.1/4.1.scm` and `cont/conts.scm` in the GitHub repository `jlam5555/sicp`. Chez Scheme 9.4 was used as the Scheme of choice.

```

(define (mi::analyze-call/ccs exp)
  ;; for (call/cc f), call f, binding the continuation
  ;; as a primitive procedure to the argument of f
  (let ([fproc (mi::analyze (mi::call/ccs-arg exp))])
    (lambda (env ccs)
      (fproc
       env
       (cons
        (lambda (proc)
          (mi::execute-application
           proc
           (list
            (map (lambda (cc) (list 'continuation cc)) ccs))
            ccs))
        (cdr ccs))))))

(define (mi::analyze-call/new-ccs exp)
  (let ([fproc (mi::analyze (mi::call/new-ccs-generator exp))])
    [body (mi::analyze (mi::call/new-ccs-body exp))])
    (lambda (env ccs)
      (fproc
       env
       (cons
        (lambda (proc)
          (mi::execute-application
           proc
           (list
            (map (lambda (cc) (list 'continuation cc)) ccs))
            (cons
             (lambda (new-ccs)
               ;; no safety checks!!! run body with
               ;; user-supplied ccs mapped to primitive procs
               (body env
                (map
                 (lambda (cc)
                  (lambda (val)
                   (mi::execute-application
                    cc
                    (list val)
                    ccs)))
                 new-ccs)))
             (cdr ccs))))
        (cdr ccs))))))

```

Figure 14: Implementations of call/ccs and call/new-ccs

```

(define (try/catch try catch)
  ;; try is a thunk
  ;; catch is a proc that takes the error msg as input
  ;; both try and catch return a value
  (call/new-ccs
   (lambda (ccs)
     (cons (car ccs)
           (cons catch
                 (cdr (cdr ccs))))))
   (try)))

(define (my-/ x y)
  ;; modified version of `/` that calls the error
  ;; continuation on div0
  (if [= y 0]
      (throw "my-/: /0")
      (/ x y)))

(try/catch
 (lambda () (my-/ 1 2))
 (lambda (err) "oh no"))
;; => 1/2

(try/catch
 (lambda () (my-/ 1 0))
 (lambda (err) "oh no"))
;; => "oh no"

(try/catch
 (lambda () (throw 32))
 (lambda (err) (throw 54) "oh no"))
;; => "top-level error continuation called with irritant 54"

```

Figure 15: `try/catch` implementation and usage

```

(define (current-continuation)
  (call/cc (lambda (cc) (cc cc))))

(define (tree-iterator tree)
  (lambda (yield)
    (define (walk tree)
      (if [not (pair? tree)]
          (yield tree)
          (begin
             (walk (car tree))
             (walk (cdr tree))))))
    (walk tree)))

(define (make-yield for-cc)
  (lambda (value)
    ;; `yield` implementation
    (define cc (current-continuation))
    (if [procedure? cc]
        ;; when called from generator, return to for loop cont
        (for-cc (cons cc value))
        ;; when called from for loop, return to generator cont
        (void))))

(define (for-generator iterator body)
  (define (loop iterator-cont)
    (define cc (current-continuation))
    (if [procedure? cc]
        ;; get next value using the generator continuation
        (if iterator-cont
            (iterator-cont (void))
            (iterator (make-yield cc)))
        ;; value handler: receive new value and continuation
        [begin
         (body (cdr cc)) ; next generator value
         (loop (car cc))] ; next generator continuation
        ;; begin iterator loop with no iterator continuation
        (loop false)))

(define (println val) (display val) (display "\n"))

(for-generator (tree-iterator (cons 3 (cons (cons 4 5) 6)))
  println)
;; => 3 4 5 6

```

Figure 16: Demonstration of generators using `call/cc`

## 4 Conclusion

I was able to explore continuations to a much greater depth of understanding than before. Perhaps more importantly, this discussion has piqued my interest on a number of other related topics, including: segmented and cactus stacks; delimited continuations; implementation of Javascript [Promises](#) and their relationship to monads; goroutine implementation and relationship to coroutines; reading Andrew Appel's *Compiling with Continuations*.

Unfortunately, my extension to the continuations API is probably not that useful; with the exception of the simple implementation of try/catch in programmer-space, I could not think of any other use cases where this API is more useful than simply passing around multiple continuations (or simply callbacks) as needed. Efficiency was not a concern during this project, but this new feature probably adds a nontrivial overhead to function calls as well.

## 5 References

Since most of the sources are not academic with auto-generated BibTeX citations, and are easily accessible via the web, I have chosen to include the sources as links throughout the document in lieu of a more formal citation style.

# Proposal for an independent study on the design of (mostly functional) languages

(a.k.a., a case for studying Scheme and *SICP*)

Jonathan Lam  
to Prof. Sable

January 15, 2021

(I make a lot of references to the various resources in the section Readings to motivate LISP and *SICP*, so feel free to check that out first.)

## 1 Summary

This is a report intended to motivate an independent study on the design of a language, using Lisp as the object of consideration due to its metaprogramming capabilities.

## 2 Motivation (and responses to concerns)

My technical experience of the past roughly eight years comprises that of a: polyglot programmer; full-stack developer; (somewhat) dev-ops intern; computer center operator; programming mentor; founder of various high-school programming clubs; freelance web developer; computer hardware hobbyist; collector of old server hardware; independent driver developer; \*nix power user; and, most recently student in deep learning and artificial intelligence. Through these roles I've become fairly convinced that the quality of one's code is paramount to how well a program performs; how performant, scalable, and maintainable the source tree is; and how easy it is for others to understand and contribute to a project. My concern is that, having viewed many peers' code via mentoring or group projects (or my coworkers' code when I was an intern), and having reflected on my own code from the past, that good coding quality standards are rarely taught nor enforced in many situations where a deadline must be met (particularly susceptible are academic projects and scientific computing). I wish for a more formal practice; to be able to think regularly in terms of better design principles and meaningfully-structured code.

The second issue is that I would like to pursue some topic in the space of operating systems, compilers, or language design: something in the infrastruc-



ture of the programming stack. I feel that my work has been saturated with an interest in the metaprocess of building programs; I always want to know how the libraries, the languages that comprise those libraries, and the systems that host those languages are built. In other words, I want to understand the series of abstractions, from low-level computing to the design of practical high-level libraries. In the end, I hope to perform research in this field for a graduate degree, but I am not sure what in this somewhat-large space on which to focus.

The intersection of these two topics is *Structure and Interpretation of Computer Programs* (*SICP*). The forward and prefaces to the book nicely declares its intent; here is an excerpt from the prologue by Alan Perlis:

Unfortunately, as programs get large and complicated, as they almost always do, the adequacy, consistency, and correctness of the specifications themselves become open to doubt, so that complete formal arguments of correctness seldom accompany large programs. Since large programs grow from small ones, it is crucial that we develop an arsenal of standard program structures of whose correctness we have become sure – we call them idioms – and learn to combine them into larger structures using organizational techniques of proven value. These techniques are treated at length in this book, and understanding them is essential to participation in the Promethean enterprise called programming. More than anything else, the uncovering and mastery of powerful organizational techniques accelerates our ability to create large, significant programs.

That is, *SICP* is a book about abstraction as a building block of building practical software. Not only that, but the last two chapters specifically focus on designing a new LISP interpreter as an example of a large useful program built with the programming practices emphasized in the earlier parts of the book. The book uses LISP as the object of focus due to its powerful metaprogramming capabilities.

My goal is to turn this into an independent study over the summer so I can have a mentor to ease me through the teachings of this book, to have someone to discuss inquiries with, and to keep me focused. The ultimate goal is that this, along with the related exercises and readings, will help me explore into the field of language design principles so that I can smoothly transition into a research project by the fall semester.

I don't think I need to elaborate on why LISP such a fascinating language, both in its history and its structure, but I want to defend the argument that it (and *SICP*) can still be the language (and book) of choice today in an academic setting.

## 2.1 Modern irrelevance – Choice of book and language

The largest concern during our meeting last semester was that both LISP and *SICP* are old, and its study may not be the most fruitful. This issue seems

critical, given that many CS programs use languages that are more modern and relevant to practical programming, such as C++, Java, and Python; even the courses that have traditionally used LISP (and SICP) to demonstrate some fundamentals of programming (most notably CS61A at Berkeley and 6.001 at MIT) have switched to Python-based curricula. Similarly, *How to Design Programs (HtDP)* was written to address some of the pedagogical concerns about *SICP*. Similarly, it is concerning that many of the reviews of *SICP* come from over a decade ago.

However, I would like to dispel these general concerns. First of all, there are innumerable sources that praise *SICP* for its timeless and language-agnostic ability to change the way a programmer thinks about his code from the building blocks available in the language. With what little experience I have in Scheme, I truly feel this is the case – while barebones LISP cannot compete with high-performance languages (e.g., FORTRAN, C, C++), domain-specific languages (e.g., SQL), or domain-specific libraries (e.g., TensorFlow), there are many appraisals of LISP as teaching you how you can build abstractions in software using the building blocks those languages and libraries have to offer you.

To argue about the relevance of *SICP*, I hope that the fact that they are still being offered at major universities (e.g., 6.037 at MIT), albeit not as the introductory course, speaks for itself. (Similarly, Berkeley’s 61A teaches Python in the spirit of *SICP*, and its official course name is still “Structure and Interpretations of Computer Programs.”) I also believe that the reasons why these prestigious schools stopped teaching with LISP in their introductory courses (see “Programming by Poking: Why MIT stopped teaching SICP”) says something about the changing nature of what it means to work in a CS-related job (i.e., now, it is about “poking” around software libraries rather than really doing the engineering from the most basic building blocks), but my goal is more similar to the original ideal: to be able to understand how to manipulate programming languages well enough from their most basic building blocks by means of abstraction. And that is exactly what *SICP* is about.

Despite its old age, LISP is also still innovating: just looking at Scheme, for example, we have the evolving standard. Currently, seven versions of the Scheme standard (currently at R6RS and R7RS) have been released. The language offers several features that have yet to become a part of any other mainstream language, such as first-class continuations or hygienic macros, and its macro system is still unrivaled among mainstream language (the only other languages I’ve heard it’s comparable to are metaprogramming languages such as MetaLua). There are several Scheme variants that are actively maintained and optimized, such as Chez Scheme by Cisco; other LISP variants, such as Common Lisp and Clojure, have major industry support (and the latter is a JVM language, and thus is as up-to-date as the JVM platform is). Performance-wise, many LISP compilers and interpreters have gotten to the point where performance loss compared to high-performance languages like C is negligible, and most variants of LISP beat out other high-level interpreted languages such as Python (at least plain Python without its BLAS or CUDA integrations). (And for many projects, the computing power of most devices is more than sufficient

for any modern language.) Along with the general fact that you can easily build abstractions upon abstractions (according to “Lisp Is Too Powerful”, it “allows people to invent their own little worlds”) means that using an efficient declarative language like LISP is a way to prevent accruing technical debt.

I know that some of these claims are weak, but the original concern that *SICP* and *LISP* are out-of-date are also based on pure speculation. With a good deal of Internet searching I was unable to find an example of a modern language feature that LISP was incapable of doing, outside of the typical pitfalls of declarative languages (e.g., explicit memory management, array-based data structures such as hashtables, static typing, etc.), and was unable to find any major modern innovation in functional programming that LISP specifically cannot perform. However, to be complete, I have linked *The Journal of Functional Programming* as a possible resource for this independent study, which could be useful for contemporary research (but this may not be very relevant to LISP or the main task at hand).

As an aside, I mentioned (*HtDP*) was a textbook intended to address some of the academic shortcomings of *SICP*, but I believe that many of those changes are actually in conflict with what I intend to achieve. While I have not read “The Structure and Interpretation of the Computer Science Curriculum” (the paper describing the differences between the books) in detail, from a quick perusal it seems that *HtDP* was intended to be a more student-friendly introduction, with a more gradual learning curve and explicit teachings. From what I understand, it also does not go into all the fuss of developing an interpreter as an ultimate goal as *SICP* does, but rather focuses on the more basic design aspects. However, since *SICP* explicitly goes into the details of building an interpreter; since I have a basic knowledge in LISP; and since I am not averse to the challenging problems (rather, I believe that struggling through these problems slowly are key to learning a topic well), I believe that the classic *SICP* will be the better textbook to achieve my goal.

## 2.2 Devil’s advocate: why *not* a more modern language?

Rather than defend LISP, I now present an opposing view: why *not* use a language like Python, JavaScript, or Scala, which are undoubtedly more relevant nowadays? What does LISP have to offer that the others don’t?

This concern was brought up in our meeting last semester, and my response was that LISP is known as the “programmable programmer’s language,” and I mumbled something about its homoiconicity, and that you could easily extend the syntax of LISP. But your rebuttal is that you can’t change the structure of LISP, any more than you can’t change the structure of another language like C or Python. Plus, if we’re talking about macros, C already has macros – what makes LISP any different or better?

I spent some time trying to clarify this via some Internet searches, many of which are in the Readings section below. My new answer to this is that the “programmable nature” of LISP is due to a combination of its homoiconic nature and its macros. The homoiconic nature means that the parsing of LISP

code is trivial due to its simple, recursive syntax that “can be summarized on a business card” (see “Lisp vs. Haskell”) and very clearly corresponds to an AST. This is as opposed to many other languages, which have a plethora of inconsistent syntaxes to indicate different kinds of operations, which gives concision at the cost of terribly-inconvenient parsing. As a result of this, many languages that do have macros (recall that macros are functions which transforms source code; in addition to introducing syntactic sugar, they can also change control flow and thus introduce control (syntactic) primitives and otherwise) only have *lexical* macros, e.g., C’s token-replacement preprocessor, which is nice for simple replacements but nothing more than that. LISP’s *syntactic* macros, on the other hand, are well-integrated into the language and can perform arbitrary transformations (i.e., can run arbitrary logic on the input code to produce the output code, which is capable of infinitely many transformations and not simply text replacement), as well as provide features such as hygienic macros that are intractable (if not impossible) with lexical macros.

The fact that LISP has these syntactic macros makes it an ideal candidate for a language to build abstractions out of. No, we can’t change the barebones syntax of LISP (i.e., parentheses, spaces, and tokens must follow the correct syntax), but we can define new syntactic primitives (e.g., control flow, text replacement, data-as-code, etc.) that allow us to extend the language. (If this is unclear, there are a number of resources in the Readings section which do a better job explaining this than I do here.) This undoubtedly was a key reason that the authors of *SICP* chose LISP, and I can see it as an excellent choice for experimenting with different language constructs and styles.

## 2.3 Relationship with functional languages?

From the video call earlier, it seemed like your impression of the book (and this independent study) was that it might be a study of functional languages – this is not true. I believe you mentioned that you might consider teaching a course on functional programming in the future at Cooper, which is an endeavor I admire as a good step towards creating better programmers – but that is not my goal here. While Lisp is (primarily) functional, that is neither its defining characteristic nor the reason why it is so powerful.

I gave a talk about Lisp last semester in an IEEEExACM-organized educational session (rough transcription) because I love the language and hope to indoctrinate as many people as I can. In the original flyer. In the event poster, I called the event “An introduction to Scheme Lisp and Functional Programming,” but when I was generating the presentation I was wondering if this was what I really wanted to convey. Functional programming indeed is something that I don’t believe many Cooper students are adequately exposed to (the main culprit being firstly that C and Python being taught in terms in data structures (which are mostly imperatively-defined, see “Disadvantages of purely functional languages”), and secondly that scientific programming doesn’t at all promote non-imperative-style coding practices), but it didn’t seem to me that it was the selling feature of Lisp. After all, many modern languages are

multi-paradigm now (e.g., modern high-level dynamic programming languages like Python, Scala, JavaScript, Ruby, and even recent editions of traditionally-imperative languages like C++ and Java), so this is not at all impressive anymore, even if LISP was the pioneer in many of the topics in functional programming. Rather, I believe that the fact that Lisp’s powerful syntax inherently follow a recursive, expression-based, and therefore declarative nature (since its inception in McCarthy’s original paper) is a byproduct of its metaprogramming capabilities, and should be treated as such. As mentioned earlier, I have included *The Journal of Functional Programming* as a possible resource for this class, which should more than suffice innovations in functional programming, but this study is orthogonal to LISP (as it should be).

For a study of functional programming, being forced to work in a purely-functional language such as Haskell or Prolog (the latter of which is a Lisp) is probably a better choice. To address the first concern about modernness, there are some resources in the Readings section that illustrate that Haskell is still relevant and performant, despite being function and high-level.

## 2.4 Lack of knowledge of LISP for mentorship

You voiced your concern that your last experience with LISP was many years ago, and may not have enough experience with the language to be able to advise me if the independent study were to focus on it. However, my view as a polyglot programmer is that that fact should never really be a problem. Also, I firmly believe that having a second person to discuss ideas with, even if they aren’t designated experts in the field, is incredibly useful. Perhaps discussions with someone who is the opposite of an expert – an inquisitive, but otherwise clueless on the matter – is the best way to learn, because then more naïve questions are posed and answered, and you simply double the brainpower. This is a fairly general and idealistic pedagogy, but I find discussions with both informed and uninformed people very useful when self-learning any material.

Of course, you (Prof. Sable) are not generally naïve when it comes to programming languages or the principles of data structures and algorithms (and I like to believe that I am also knowledgeable to some degree), so and this peripheral knowledge should accelerate the rate of learning.

## 2.5 CL vs. Scheme (vs. Others)

This is related to the concerns about the modernness and usefulness of Scheme, which is such a minimal language that it doesn’t have many claims to fame outside of academia (and principally, outside of the realm of *SICP*). I remember you suggested that, instead of Scheme, we could use Common Lisp, which has many more libraries and is thus more “practical.” However, I would like to argue this point on three grounds:

Firstly, that there is more than enough mental material to fill a semester (or two, or three) with purely academic exercises in Scheme, using only the exercises in *SICP* or additional exercises (e.g., from *The [Little/Intermediate/Seasoned]*

*Schemer*). Satisfactorily learning Common Lisp, a much more featured language with larger industry support, very well might be overwhelming for a semester

Secondly, that the ideas in *SICP* are largely language-agnostic (given that other general-purposes languages are less-capable of syntactic abstraction but more featured with their default syntax and libraries). Of course, modern vocational programming or scientific computing is highly language- and library-specific, and doesn't require the kind of analytical, problem-solving skillset that *SICP* offers. Perhaps that makes Scheme less appealing to students who simply want a quick job with their EE/CS degree, but again, that is not the goal for me here.

Thirdly, the fact that Scheme has not gained widespread appeal and is not useful for practical coding is probably largely true, but I have personally found it to be useful and somewhat widespread. In my casual programming experience, I have seen Scheme (mostly Guile Scheme due to its GNU integration) used in various useful projects, such as for xbindkeys (keyboard-to-macro mappings in Linux) and emacs lisp. Similarly, mainstream learning resources and coding platforms for Scheme is not at all limited: an excellent tool I have used in the past to learn new languages, exercism.io, includes a track for (Chez and Guile) Scheme. And for school and hobby projects (except perhaps for high-performance computing, such as deep learning and scientific computing that needs intimate control over memory), I've found Scheme to be plenty useful, e.g., for use in ECE469 Artificial Intelligence and MA352 Discrete Mathematics. I'm sure that I would have used Scheme in many other courses had I been familiar with it earlier.

## 2.6 Overlap with ECE466 (Compilers)

Knowing that *SICP* builds up towards writing a Lisp interpreter might make it seem to have significant overlap with a compilers course. However, a cursory look at the ECE466 course webpage indicates that the topics covered by this course are very different than the goals of *SICP*. While the compilers class shows how to build a “practical compiler,” e.g., ASTs, parsing, control flow optimizations, etc., *SICP* is designed to illustrate how to build more complex applications by abstracting the building blocks of a fully-featured low-level language. In other words, it is not the end goal of Lisp to build a programming language compiler or interpreter, but it is a very illustrative example that shows the power of metaprogramming; this is reinforced by the view in “Short explanation of last two chapters in *SICP*.” Furthermore, the methods with which *SICP* attempts to build a interpreter are fundamentally different from those of compilers, as the building blocks are very different: parsing and memory management are already implicit in Lisp, whereas any practical compiler design is lower-level and would need to take this into consideration. All in all, I hope that both ECE466 and this study will aid in selecting a research topic.

### 3 (Preliminary) (Tentative) Timeline

I believe that following the book should offer more than enough material for a course semester. As many testimonies of the book say, *SICP* is a tome that takes time to read and digest, and its exercises are not trivial. Even at prestigious universities such as Berkeley and MIT (at MIT it is taught by the original authors of the book), this is offered as a semester-long course (and the authors state in the forward that it is more material than what can be covered in a typical semester), and is used as the textbook for graduate-level courses on symbolic programming as well (see MIT's 6.945 in the Materials). We can follow a schedule similar to that posted on MIT or Berkeley's course webpages, and use the video lectures from MIT. Extra exercises may be pulled from *The [Little/Intermediate/Seasoned] Schemer*.

However, to satisfy the need to feel up-to-date, perhaps additional readings should be completed every week (in the tradition of Cooper's ECE472 Deep Learning course, perhaps a set of 2-5 weekly readings from *The Journal of Functional Programming*), or some of the other scholarly articles or books on Lisp listed below, may serve as additional reading.

I have already covered the first chapter of the book on my own, and have a working knowledge of the basics of LISP through school projects, so I believe we can dive straight into the second chapter. If we plan to cover the latter four chapters of the book for a roughly 13-week agenda for the summer semester, a possible weekly agenda for the independent study might look like:

1. §2.1-2.4: Symbolic data, data abstraction, data hierarchy and closures
2. §2.5: Systems with generic operations
3. Continuation of §2.5: Implement a symbolic algebra system (perhaps with macros and symbolic differentiation?)
4. §3.1-3.2: Local state and the evaluation model
5. §3.3-3.4: Mutable data and concurrency
6. §3.5: Streams (and perhaps a project to implement streams)
7. §4.1: The metacircular evaluator
8. §4.2-4.3: Lazy evaluation, nondeterministic computing
9. §4.4: Logic programming: (perhaps a project to implement a Prolog-like language)
10. §5.1-5.2: Register machines
11. §5.3: Allocation and garbage collection
12. §5.3: The explicit-control evaluator
13. §5.4: Compilation

## 4 (Preliminary) (Tentative) Deliverables and Assessment

Assignments from *SICP* and *The [Little/Intermediate/Seasoned] Schemer* could be used as weekly assignments, as well as reading reports. The final project may be a Lisp interpreter written in Lisp, following the trajectory of the book.

## 5 Materials

- *Structure and Interpretation of Computer Programs (SICP)*  
a.k.a. the “Wizard book,” and the subject of discussion of most of this article
- Structure and Interpretation of Computer Programs (University of California, Berkeley CS61A)  
Python in the tradition of *SICP* (Associated materials: Composing Programs)
- Structure and Interpretation of Computer Programs (MIT 6.001)  
Original course taught by the authors of the book (stopped being taught in early 2000’s)
- Structure and Interpretation of Computer Programs (MIT 6.037)  
Newer course taught by the authors of the book (2019)
- Adventures in Advanced Symbolic Programming (MIT 6.945)  
Graduate course on symbolic programming, uses *SICP* as its textbook
- Journal of Functional Programming  
If we want to stay up to date, can choose a few articles out of this every week. Unfortunately may need to pay to subscribe
- *On Lisp*  
A whole book on Lisp by Paul Graham (who is described in the article “How Lisp became God’s Own Programming Language” below)
- *The Common Lisp Cookbook*  
The title is self-explanatory
- R6RS  
Latest widely-adopted Scheme standard, including standard libraries
- Recursive Functions of Symbolic Expressions and Their Computation by Machine, Part I  
Original paper by McCarthy on S-expressions, which formed the basis for Lisp; Lisp was implemented shortly after this paper using these principles



- *The [Little/Intermediate/Seasoned] Schemer*  
A series of problems, separated by difficulty, designed to illustrate various aspects of the Scheme language
- *Let over Lambda: 50 years of Lisp*  
A guidebook to Common Lisp, including more advanced features that you wouldn't normally find elsewhere; first six chapters are available online for free
- *How to Design Programs* (HtDP)  
A book mostly in the tradition of *SICP*, but specifically aimed at its pedagogical shortcomings (e.g., HtDP is intended to require less domain knowledge and be more explicit in its principles); *The Structure and Interpretation of the Computer Science Curriculum* addresses these pedagogical differences. Second edition came out last year (2018)
- *Paradigms of Artificial Intelligence Programming: Case Studies in Common Lisp* (PAIP)  
Chapters include logic programming and other practical problems in Lisp (e.g., a general purpose solver (GPS), ELIZA, a computer algebra system (CAS) and other symbolic mathematics)

## 6 Readings to motivate LISP and *SICP*

I realize I skimmed quite a few Q&A posts, articles, books, etc. related to “Why Lisp?” and especially “Why Lisp macros (as opposed to macros in other languages or ordinary first-class functions)?” because macros are what really make Lisp so useful. Here is a brief summary of some articles.

- “Lisp: A language for stratified design”  
1987 research paper by the authors of *SICP* about the usefulness of Lisp as a way to show a complete view on programming techniques through abstraction
- “Scheme vs. Python”  
The title is misleading; this is written by the professor for 61A at Berkeley, which uses *SICP*, and it defends the use of the book and Scheme rather than Python. (Since this post was written, I believe 61A has switched to Python (most likely due to a reason similar to “Programming by Poking: Why MIT stopped teaching SICP”), but it still attempts to follow the tradition of *SICP*)
- “Is there an expiration date for well regarded, but old books on programming?”

No, most of the concepts in these books are fundamental (and therefore timeless), just as the basics of mathematics (e.g., calculus) hasn't changed for centuries; also many concepts are language-agnostic

- “Programming by Poking: Why MIT stopped teaching SICP”

(The title is misleading: MIT still has courses that use and teach *SICP*, but the iconic 6.001 course that was the introduction to all EECS undergrads at MIT was abolished in favor of an introductory course in Python, 6.0001)

Unfortunately, much of the real work programmers do nowadays is “programming by poking”: messing around with pieces of code (high-level languages and libraries) to just get it to work, people don't have to build systems from the bottom up nowadays, so they switched the introductory course to python

Personally, I find this reason very sad and reflects the state of most programmers nowadays. But this is contradictory to my motivation (see section above)

- “Notes on Structure and Interpretation of Computer Programs”

Mentions a few notable topics that *SICP* teaches: homoiconicity, how to write an interpreter, register machines, etc.

- “Why is Haskell (GHC) so darn fast?”

A discussion on the speed of modern functional languages, and how they achieve their impressive real-world speed (as opposed to the common mode of thinking that functional languages are typically more inefficient)

- “Disadvantages of purely functional languages”

Written by a pronounced naysayer of Haskell; mostly describes the fact that some common data structures cannot be implemented in a functional way

- “Lisp vs. Haskell”

Speaks to Lisp's metaprogramming abilities and simplicity ((almost) everything is built out of seven core functions)

- “What is so great about Lisp?”

General discussion on why so many people admire Lisp. Mostly emphasizes on the simple syntax, which leads naturally to a metacircular interpreter, as well as bringing up a host of interesting readings and Green-spun's tenth rule

Also mentions a few disadvantages, such as the relatively small number of libraries, difficulty to pick up, and being forced to use dynamic typing (whether this is a boon or poison depends on the application). However, none of these affect learning the fundamentals in an academic context

- “Lisp Is Too Powerful”

The first sentence is illustrative: “I think one of the problems with Lisp is that it is too powerful. It has so much meta ability that it allows people to invent their own little worlds, and it takes a while to figure out each person’s little world”

- “What makes Lisp a programmable programming language? Why does this matter? What are the advantages of being one?”

Describes what it means when people describe Lisp as a “programmable programming language” – namely, that it can

Also describes how Haskell doesn’t achieve these goals, which shows that (even purely) functional programming does not amount to the metaprogramming capability that Lisp offers. Is probably more true for less-functional languages like Python

- “How Lisp became God’s Own Programming Language”

Perhaps my favorite articles on “Why Lisp” that feels complete in multiple ways. Goes into a history of Lisp (including its motivations and impacts). The last paragraph, about the revival of Lisp after the AI winter and its lingering impact on “modern” languages like Ruby and Python, is especially satisfying.

- “What makes Lisp macros so special?”

Many good answers here. E.g., can implement list comprehension or infix notation in Lisp, write code replacements (which is more advanced than the C’s token-expansion, write something like C#’s LINQ notation; a macro is a function that takes in some source code and outputs source code, but it can perform logic on the expansion as well, including processing the arguments to the macro), changing execution order, the fact that Lisp’s homoiconic nature makes macros much more useful than other languages with macros (“in C, you would have to write a custom pre-processor [which would probably qualify as a sufficiently complicated C program]”)

- “What are Lisp macros good for, anyway?”

Illustrative example of a timing macro in Lisp, and comparison to purely functional approaches (i.e., without macros) in Java and Python. Illustrates how macros allow for a different evaluation pattern for its arguments, i.e., rewriting the evaluation path, which cannot be done simply with functions (i.e., we can create new control primitives just like builtin control flow operators)

- “What can you do with Lisp macros that you can’t do with first-class functions?”

Relevant to the previous resource. Useful comment on an answer: “To abstract and hopefully clarify a bit, Lisp macros allow you to control

whether (and when) arguments are evaluated; functions do not. With a function, the arguments are always evaluated when the function is invoked. With a macro, the arguments are passed to the macro as ordinary data. The macro code then decides itself if and when they should be evaluated. This is why you can't write `reverse_function` without macros: as a function, the argument would be evaluated (and cause an error) before the function body gets a chance to rewrite it."

This also goes over some of the basic ideas, e.g., that macros do expansion at compile-time whereas first-class functions create closures at runtime and have the extra overhead of a function call; this has the added benefit of greater efficiency (similar to macros in any other language)

- Paul Graham's essays (mostly on programming languages)

Haven't read these yet, but a good number are on Lisp. Paul Graham is a co-founder of ycombinator and helped to revive Lisp after the AI winter with these essays (for more history, see "How Lisp became God's Own Programming Language")

- "Automata via Macros"

Research paper aiming to bring light to the relevance and need of Lisp-like macros, since they are not widely used in more modern languages (e.g., Java and C). Also focuses on the need for tail-call optimization for macros to be useful and efficient.

Also has a very succinct list of overview of uses of macros: "providing cosmetics; introducing binding constructs; implementing 'control operators', i.e., ones that alter the order of evaluation; defining data languages"

- "Why do people say Lisp has true macros versus C/C++ macros?"

A question dedicated to the comparison of Lisp-style macros to those that exist in other languages. The main point is that since the language is homoiconic, you have the power of the entire Lisp language at your hands to transform the input code to the output code, whereas in C/C++ the language is much more complicated and only does token replacements. Similarly, the C preprocessor is very decoupled from the rest of the language and doesn't have any logical utilities (other than token replacement), whereas the Lisp macro system is very intimately tied into its language, as you can perform arbitrarily-complex Lisp code to perform the code transformation

- "Lisp Macro"

Similar arguments as in the previous articles, with some personal testimonies. Includes a lot of Q&A about use cases (including when not to use macros)

- "Why aren't macros included in most modern programming languages?"

Yet another explanation of macros and how they are different in C and Lisp. While C macros are lexical and carry no syntactic meaning, Lisp macros are syntactic and can be treated as regular code. (E.g., one implication is the use of hygienic macros in Lisp, which safely do not shadow variables in the outside scope)

- “What makes macros possible in lisp dialects but not in other languages?”

An answer to this question with some technical explanation

- “Short explanation of the last two chapters of SICP”

Explains that the final two chapters are not a repeat of a compilers course, but show that an interpreter is just a computer program that takes data as input, and thus is a good example of data abstractions

- “SICP - Worth it?”

Relevant quote: “It’s a subtle book. As others say, you do things in it you probably won’t ever do again, in a language you’ll probably never use again... But it can change (and improve) the way you write and understand programs, in any language.”